

# Databases

Gaurav Sood

Spring 2015

Data

# Data

- Bits on non-volatile storage e.g. magnetic media, SSD

# Data

- Bits on non-volatile storage e.g. magnetic media, SSD
- Not just bits on hard disk, but organized bits (a data model)

# Data

- Bits on non-volatile storage e.g. magnetic media, SSD
- Not just bits on hard disk, but organized bits (a data model)
  - Nested folders  $\sim$  tree

# Data

- Bits on non-volatile storage e.g. magnetic media, SSD
- Not just bits on hard disk, but organized bits (a data model)
  - Nested folders  $\sim$  tree
  - Rows and columns  $\sim$  table

## Data Model

- Structure

# Data Model

- Structure
  - Rows and columns (Relational Database)



# Data Model

- Structure
  - Rows and columns (Relational Database)
  - Key-value pairs (NoSQL systems)

# Data Model

- Structure
  - Rows and columns (Relational Database)
  - Key-value pairs (NoSQL systems)
  - Sequence of bytes (Files)

# Data Model

- Structure
  - Rows and columns (Relational Database)
  - Key-value pairs (NoSQL systems)
  - Sequence of bytes (Files)
- Constraints

# Data Model

- Structure
  - Rows and columns (Relational Database)
  - Key-value pairs (NoSQL systems)
  - Sequence of bytes (Files)
- Constraints
  - All columns must be of same length

# Data Model

- Structure
  - Rows and columns (Relational Database)
  - Key-value pairs (NoSQL systems)
  - Sequence of bytes (Files)
- Constraints
  - All columns must be of same length
  - Each value in this column has to be X

# Data Model

- Structure

- Rows and columns (Relational Database)
- Key-value pairs (NoSQL systems)
- Sequence of bytes (Files)

- Constraints

- All columns must be of same length
- Each value in this column has to be X
- Child cannot have two parents (tree) (before Gmail Labels)

# Data Model

- Structure
  - Rows and columns (Relational Database)
  - Key-value pairs (NoSQL systems)
  - Sequence of bytes (Files)
- Constraints
  - All columns must be of same length
  - Each value in this column has to be X
  - Child cannot have two parents (tree) (before Gmail Labels)
- (Efficiently Supported) Operations

# Data Model

- Structure
  - Rows and columns (Relational Database)
  - Key-value pairs (NoSQL systems)
  - Sequence of bytes (Files)
- Constraints
  - All columns must be of same length
  - Each value in this column has to be X
  - Child cannot have two parents (tree) (before Gmail Labels)
- (Efficiently Supported) Operations
  - Look up value of key x.



# Data Model

- Structure
  - Rows and columns (Relational Database)
  - Key-value pairs (NoSQL systems)
  - Sequence of bytes (Files)
- Constraints
  - All columns must be of same length
  - Each value in this column has to be X
  - Child cannot have two parents (tree) (before Gmail Labels)
- (Efficiently Supported) Operations
  - Look up value of key x.
  - Find me rows where column 'lastname' is Jordan.

# Data Model

- Structure
  - Rows and columns (Relational Database)
  - Key-value pairs (NoSQL systems)
  - Sequence of bytes (Files)
- Constraints
  - All columns must be of same length
  - Each value in this column has to be X
  - Child cannot have two parents (tree) (before Gmail Labels)
- (Efficiently Supported) Operations
  - Look up value of key  $x$ .
  - Find me rows where column 'lastname' is Jordan.
  - Get next  $n$  bytes.

## Questions to consider

- What are the features of the data? How do we plan to use the data?

## Questions to consider

- What are the features of the data? How do we plan to use the data?
- How often are the data updated?

## Questions to consider

- What are the features of the data? How do we plan to use the data?
- How often are the data updated?
- What queries are efficiently supported?

## Questions to consider

- What are the features of the data? How do we plan to use the data?
- How often are the data updated?
- What queries are efficiently supported?
- How are the data stored on disk – the physical data model?

## Questions to consider

- What are the features of the data? How do we plan to use the data?
- How often are the data updated?
- What queries are efficiently supported?
- How are the data stored on disk – the physical data model?
- How are data organized internally – the abstract data model?

## What are databases?

- Organized collection of data.



## What are databases?

- Organized collection of data.
- Organized to afford efficient retrieval (aim can vary).

## What are databases?

- Organized collection of data.
- Organized to afford efficient retrieval (aim can vary).
- Data + Metadata about structure and organization.

## What are databases?

- Organized collection of data.
- Organized to afford efficient retrieval (aim can vary).
- Data + Metadata about structure and organization.
- Data that are self-describing  $\sim$  have a schema.

## What are databases?

- Organized collection of data.
- Organized to afford efficient retrieval (aim can vary).
- Data + Metadata about structure and organization.
- Data that are self-describing  $\sim$  have a schema.
- Generally accessed through a Database Management System.

## What do Database Systems (DBMS) do?

- Provide efficient, reliable, convenient and safe, multi-user storage of and access to massive amounts of data.

# What do Database Systems (DBMS) do?

- Provide **efficient**, **reliable**, **convenient** and **safe**, **multi-user** storage of and access to **massive** amounts of data.
  - **Massive:** Terabytes. Handle data that reside outside memory.

# What do Database Systems (DBMS) do?

- Provide **efficient**, **reliable**, **convenient** and **safe**, **multi-user** storage of and access to **massive** amounts of data.
  - **Massive:** Terabytes. Handle data that reside outside memory.
  - **Safe:** Robust to power outages, node failures etc.

# What do Database Systems (DBMS) do?

- Provide **efficient, reliable, convenient** and **safe, multi-user** storage of and access to **massive** amounts of data.
  - **Massive:** Terabytes. Handle data that reside outside memory.
  - **Safe:** Robust to power outages, node failures etc.
  - **Multi-user:** Concurrency control. Not one user/ one user per data item. File system concurrency.



# What do Database Systems (DBMS) do?

- Provide **efficient, reliable, convenient** and **safe, multi-user** storage of and access to **massive** amounts of data.
  - **Massive:** Terabytes. Handle data that reside outside memory.
  - **Safe:** Robust to power outages, node failures etc.
  - **Multi-user:** Concurrency control. Not one user/ one user per data item. File system concurrency.
  - **Convenient:** High level query languages; declarative: what you want, not describe algorithm

# What do Database Systems (DBMS) do?

- Provide **efficient**, **reliable**, **convenient** and **safe**, **multi-user** storage of and access to **massive** amounts of data.
  - **Massive:** Terabytes. Handle data that reside outside memory.
  - **Safe:** Robust to power outages, node failures etc.
  - **Multi-user:** Concurrency control. Not one user/ one user per data item. File system concurrency.
  - **Convenient:** High level query languages; declarative: what you want, not describe algorithm
  - **Efficient:** Performance, performance, performance.

# What do Database Systems (DBMS) do?

- Provide **efficient, reliable, convenient** and **safe, multi-user** storage of and access to **massive** amounts of data.
  - **Massive:** Terabytes. Handle data that reside outside memory.
  - **Safe:** Robust to power outages, node failures etc.
  - **Multi-user:** Concurrency control. Not one user/ one user per data item. File system concurrency.
  - **Convenient:** High level query languages; declarative: what you want, not describe algorithm
  - **Efficient:** Performance, performance, performance.
  - **Reliable:** 99.99999% up time.

## Key concepts

- Data model  
XML, relational model, etc.

## Key concepts

- Data model  
XML, relational model, etc.
- Schema versus data distinction  
Types versus variables  
Schema (types, structure), actual data

## Key concepts

- Data model  
XML, relational model, etc.
- Schema versus data distinction  
Types versus variables  
Schema (types, structure), actual data
- Data definition language (DDL)  
For setting up schema

## Key concepts

- Data model  
XML, relational model, etc.
- Schema versus data distinction  
Types versus variables  
Schema (types, structure), actual data
- Data definition language (DDL)  
For setting up schema
- Data manipulation or query language (DML)  
For querying and modifying the database

## Some Data Models



# XML

- Extensive Markup Language

Similar to HTML

Tags describe data rather than how to format

# XML

- Extensive Markup Language

Similar to HTML

Tags describe data rather than how to format

- Standard for data exchange and representation

# XML

- Extensive Markup Language

Similar to HTML

Tags describe data rather than how to format

- Standard for data exchange and representation
- XML

# XML

- Extensive Markup Language

Similar to HTML

Tags describe data rather than how to format

- Standard for data exchange and representation

- XML

- Tagged elements

# XML

- Extensive Markup Language

Similar to HTML

Tags describe data rather than how to format

- Standard for data exchange and representation

- XML

- Tagged elements
  - Nesting of elements

# XML

- Extensive Markup Language

Similar to HTML

Tags describe data rather than how to format

- Standard for data exchange and representation

- XML

- Tagged elements
  - Nesting of elements
  - Attributes

# XML

- Extensive Markup Language

Similar to HTML

Tags describe data rather than how to format

- Standard for data exchange and representation

- XML

- Tagged elements
- Nesting of elements
- Attributes
- Not all elements have attributes

# Plain Text

June 5, 2006

Floor Statement of Senator Barack Obama Federal  
Marriage Amendment

I agree with most Americans, with Democrats and  
Republicans, with Vice President Cheney, with over  
2,000 religious leaders of all different beliefs,  
that decisions about marriage, as they always have,  
should be left to the states.



# XML

<DOC>

<DOCNO>obama-2006-06-05-01< /DOCNO>

<TEXT>

I agree with most Americans, with Democrats and Republicans, with Vice President Cheney, with over 2,000 religious leaders of all different beliefs, that decisions about marriage, as they always have, should be left to the states.

< /TEXT>

< /DOC>

# JSON

```
{u'category': u'Government official', u'username':  
u'RepJohnLewis', u'about': u'Official Congressional  
page for Rep. John Lewis', ... u'name': u'John  
Lewis', u'hometown': u'Troy, AL', ... u'website':  
u'johnlewis.house.gov', u'phone': u'(404) 659 -  
0116', u'birthday': u'02/21/1940', u'likes': 62974,  
...}
```

## Relational Model

- RDBMS– Codd 1970

## Relational Model

- RDBMS– Codd 1970
- Database: set of named relations (or tables)

## Relational Model

- RDBMS– Codd 1970
- Database: set of named relations (or tables)
- Each relation has a set of named attributes (or columns)

## Relational Model

- RDBMS– Codd 1970
- Database: set of named relations (or tables)
- Each relation has a set of named attributes (or columns)
- Each attribute (column) has a type (or domain)

There is also concept of an enumerated domain

## Relational Model

- RDBMS– Codd 1970
- Database: set of named relations (or tables)
- Each relation has a set of named attributes (or columns)
- Each attribute (column) has a type (or domain)
  - There is also concept of an enumerated domain
- Each row (= tuple)

## Relational Model

- RDBMS– Codd 1970
- Database: set of named relations (or tables)
- Each relation has a set of named attributes (or columns)
- Each attribute (column) has a type (or domain)  
There is also concept of an enumerated domain
- Each row (= tuple)
- Schema– structure, name, type of attributes



## Relational Model

- RDBMS– Codd 1970
- Database: set of named relations (or tables)
- Each relation has a set of named attributes (or columns)
- Each attribute (column) has a type (or domain)  
There is also concept of an enumerated domain
- Each row (= tuple)
- Schema– structure, name, type of attributes
- Everything is a table

## Relational Model

- Relationships are implicit; no pointers

# Relational Model

- Relationships are implicit; no pointers
  - Physical Data independence

# Relational Model

- Relationships are implicit; no pointers
  - Physical Data independence
  - Compare to Network Data Model etc.

# Relational Model

- Relationships are implicit; no pointers
  - Physical Data independence
  - Compare to Network Data Model etc.
- Example:

# Relational Model

- Relationships are implicit; no pointers
  - Physical Data independence
  - Compare to Network Data Model etc.
- Example:
  - [course, Student id ], [student id, student name]

# Relational Model

- Relationships are implicit; no pointers
  - Physical Data independence
  - Compare to Network Data Model etc.
- Example:
  - [course, Student id ], [student id, student name]
  - Just a shared id

# Relational Model

- Relationships are implicit; no pointers
  - Physical Data independence
  - Compare to Network Data Model etc.
- Example:
  - [course, Student id ], [student id, student name]
  - Just a shared id
  - Really bad for performance, all student names for a course, look up all the ids



# Relational Model

- Relationships are implicit; no pointers
  - Physical Data independence
  - Compare to Network Data Model etc.
- Example:
  - [course, Student id ], [student id, student name]
  - Just a shared id
  - Really bad for performance, all student names for a course, look up all the ids
  - Hierarchical can be quicker if all students in one course

# Relational Model

- Relationships are implicit; no pointers
  - Physical Data independence
  - Compare to Network Data Model etc.
- Example:
  - [course, Student id ], [student id, student name]
  - Just a shared id
  - Really bad for performance, all student names for a course, look up all the ids
  - Hierarchical can be quicker if all students in one course
  - But to look up all courses student X has taken, just as good (bad)

# Relational Vs. XML

	Relational	XML
Structure	Tables	Hierarchical
Schema	Fixed in advance	Flexible, self-describing
Queries	Easy	Querying not as easy
Ordering	Unordered	Implied order (document)

## Algebra of Tables

- In algebra, concept of algebraic optimization – Rules to simplify calculation

## Algebra of Tables

- In algebra, concept of algebraic optimization – Rules to simplify calculation
- SQL uses relational algebra

## Algebra of Tables

- In algebra, concept of algebraic optimization – Rules to simplify calculation
- SQL uses relational algebra
- All RDBMS optimize relational algebra 'calculations'

## Algebra of Tables

- In algebra, concept of algebraic optimization – Rules to simplify calculation
- SQL uses relational algebra
- All RDBMS optimize relational algebra 'calculations'
- Cost-based optimization

## Algebra of Tables

- In algebra, concept of algebraic optimization – Rules to simplify calculation
- SQL uses relational algebra
- All RDBMS optimize relational algebra 'calculations'
- Cost-based optimization
- We can leave this optimization to program



# Algebra of Tables

- Basic Relational Algebra

# Algebra of Tables

- Basic Relational Algebra
  - Union, Intersection, Difference

# Algebra of Tables

- Basic Relational Algebra
  - Union, Intersection, Difference
  - Selection,  $\sigma$

# Algebra of Tables

- Basic Relational Algebra
  - Union, Intersection, Difference
  - Selection,  $\sigma$
  - Projection,  $\Pi$

# Algebra of Tables

- Basic Relational Algebra
  - Union, Intersection, Difference
  - Selection,  $\sigma$
  - Projection,  $\Pi$
  - Join,  $\bowtie$

# Algebra of Tables

- Basic Relational Algebra
  - Union, Intersection, Difference
  - Selection,  $\sigma$
  - Projection,  $\Pi$
  - Join,  $\bowtie$
- Extended Relational Algebra:

# Algebra of Tables

- Basic Relational Algebra
  - Union, Intersection, Difference
  - Selection,  $\sigma$
  - Projection,  $\Pi$
  - Join,  $\bowtie$
- Extended Relational Algebra:
  - Duplicate elimination,  $d$

# Algebra of Tables

- Basic Relational Algebra
  - Union, Intersection, Difference
  - Selection,  $\sigma$
  - Projection,  $\Pi$
  - Join,  $\bowtie$
- Extended Relational Algebra:
  - Duplicate elimination,  $d$
  - Grouping and aggregation,  $g$



# Algebra of Tables

- Basic Relational Algebra
  - Union, Intersection, Difference
  - Selection,  $\sigma$
  - Projection,  $\Pi$
  - Join,  $\bowtie$
- Extended Relational Algebra:
  - Duplicate elimination,  $d$
  - Grouping and aggregation,  $g$
  - Sorting,  $t$

# Algebra of Tables

- Basic Relational Algebra
  - Union, Intersection, Difference
  - Selection,  $\sigma$
  - Projection,  $\Pi$
  - Join,  $\bowtie$
- Extended Relational Algebra:
  - Duplicate elimination,  $d$
  - Grouping and aggregation,  $g$
  - Sorting,  $t$
- Sets Vs Bags

# Algebra of Tables

- Basic Relational Algebra

- Union, Intersection, Difference
- Selection,  $\sigma$
- Projection,  $\Pi$
- Join,  $\bowtie$

- Extended Relational Algebra:

- Duplicate elimination,  $d$
- Grouping and aggregation,  $g$
- Sorting,  $t$

- Sets Vs Bags

- Sets:  $\{a,b,c\}$  (RA, Papers)

# Algebra of Tables

- Basic Relational Algebra

- Union, Intersection, Difference
- Selection,  $\sigma$
- Projection,  $\Pi$
- Join,  $\bowtie$

- Extended Relational Algebra:

- Duplicate elimination,  $d$
- Grouping and aggregation,  $g$
- Sorting,  $t$

- Sets Vs Bags

- Sets:  $\{a,b,c\}$  (RA, Papers)
- Bags:  $\{a,a,b,c\}$  (SQL)

## Select Rows

- Tuples (rows) that satisfy a particular criteria
- $\sigma_{\text{condition}}$  Relation  
Pick Certain Rows

## Select Rows

- Tuples (rows) that satisfy a particular criteria
- $\sigma_{\text{condition}}$  Relation  
Pick Certain Rows
- $\sigma_{GPA > 3}$  Student  
 $\sigma_{GPA > 3 \wedge LastName == 'Smith'}$  Student

## Select Rows

- Tuples (rows) that satisfy a particular criteria
- $\sigma_{\text{condition}}$  Relation  
Pick Certain Rows
- $\sigma_{GPA>3}$  Student  
 $\sigma_{GPA>3 \wedge LastName == 'Smith'}$  Student
- In SQL, selection using 'Where' (and project using 'Select')

# Project

- $\Pi_{A_1, A_2}$  Relation  
Pick Certain Columns



# Project

- $\Pi_{A_1, A_2}$ Relation  
Pick Certain Columns
- $\Pi_{SSN, GPA}$ Student

## Project

- $\Pi_{A_1, A_2}$ Relation  
Pick Certain Columns
- $\Pi_{SSN, GPA}$ Student
- Both Select and Project  
 $\Pi_{SSN, GPA}(\sigma_{GPA > 3} \text{Student})$   
 $\Pi_{A_1, A_2}$ Expression  
 $\sigma_{\text{condition}}$ Expression

# SQL

```
Select Col1, Col2 (or *)  
FROM R1  
Where Condition
```

## Cross-Product

- Cartesian Product

## Cross-Product

- Cartesian Product
- $R1 \times R2$ ; Every tuple in  $R1$  with each tuple in  $R2$

## Cross-Product

- Cartesian Product
- $R1 \times R2$ ; Every tuple in  $R1$  with each tuple in  $R2$
- Size of output:  $R1 * R2$

## Cross-Product

- Cartesian Product
- $R1 \times R2$ ; Every tuple in  $R1$  with each tuple in  $R2$
- Size of output:  $R1 * R2$
- `SELECT *`  
`FROM R1 CROSS JOIN R2;`

## Cross-Product

- Cartesian Product
- $R1 \times R2$ ; Every tuple in  $R1$  with each tuple in  $R2$
- Size of output:  $R1 * R2$
- `SELECT *`  
`FROM R1 CROSS JOIN R2;`
- `SELECT *`  
`FROM R1, R2;`



## Natural Join

- Means Equi-Join

## Natural Join

- Means Equi-Join
- For every record in  $R_1$ , find a tuple in  $R_2$  that satisfies a particular condition

## Natural Join

- Means Equi-Join
- For every record in R1, find a tuple in R2 that satisfies a particular condition
- Select \*  
From R1, R2  
Where  $R1.A = R2.A$

## Natural Join

- Means Equi-Join
- For every record in R1, find a tuple in R2 that satisfies a particular condition
- Select \*  
From R1, R2  
Where  $R1.A = R2.A$
- Or, do the cross-product and then filter

## Natural Join

- Means Equi-Join
- For every record in R1, find a tuple in R2 that satisfies a particular condition
- Select \*  
From R1, R2  
Where  $R1.A = R2.A$
- Or, do the cross-product and then filter
- Select \*  
From R1 Join R2  
On  $R1.A = R2.B$

## Theta Join

- Join without the equality condition

## Theta Join

- Join without the equality condition
- Equi-Join is a special case of theta join

## Theta Join

- Join without the equality condition
- Equi-Join is a special case of theta join
- Find ALL weather stations within 5 miles of centroid of a zip code



## Theta Join

- Join without the equality condition
- Equi-Join is a special case of theta join
- Find ALL weather stations within 5 miles of centroid of a zip code
- Select Distinct WStations  
From Wstations h, ZipCode s  
Where  $\text{distance}(h.\text{location}, s.\text{location}) < 5$
- ‘distance’ is a user-defined function

## Theta Join

- Join without the equality condition
- Equi-Join is a special case of theta join
- Find ALL weather stations within 5 miles of centroid of a zip code
- Select Distinct WStations  
From Wstations h, ZipCode s  
Where  $\text{distance}(h.\text{location}, s.\text{location}) < 5$
- ‘distance’ is a user-defined function
- Band Joins or Range Joins

## Theta Join

- Outer Join

## Theta Join

- Outer Join
- All tuples in R1 and pad out other values with NULLs

## Theta Join

- Outer Join
  - All tuples in R1 and pad out other values with NULLs
  - Left outer join
- ```
Select *  
From R1 LEFT OUTER Join R2  
On R1.A = R2.B
```

## Theta Join

- Outer Join
- All tuples in R1 and pad out other values with NULLs
- Left outer join  
Select \*  
From R1 LEFT OUTER Join R2  
On R1.A = R2.B
- Right outer join

## Theta Join

- Outer Join
- All tuples in R1 and pad out other values with NULLs
- Left outer join  
Select \*  
From R1 LEFT OUTER Join R2  
On R1.A = R2.B
- Right outer join
- Full outer join

## Union

- $R1 \cup R2$

Removes duplicates by default



## Union

- $R1 \cup R2$   
Removes duplicates by default
- `Select * FROM R1`  
`UNION`  
`Select * FROM R2`

## Union

- $R1 \cup R2$   
Removes duplicates by default
- `Select * FROM R1`  
`UNION`  
`Select * FROM R2`
- If you wanted duplicates, you want to do  
'UNION ALL'

Difference

$$- R1 - R2$$

## Difference

- $R1 - R2$
- Select \* From R1  
Except  
Select \* From R2

## Difference

- $R1 - R2$
- Select \* From R1  
Except  
Select \* From R2
- Looks up all tuples in R1 and takes out tuples that it also sees in R2  
R2 can have other tuples  
Everything in R1, remove things that also appear in R2

# Intersection

- Not a fundamental operator

## Intersection

- Not a fundamental operator
- $R1 \text{ Intersect } R2 = R1 - (R1 - R2)$

## Intersection

- Not a fundamental operator
- $R1 \text{ Intersect } R2 = R1 - (R1 - R2)$
- Can also express it as join



## (Virtual) Views

- Defining and Using (Virtual) Views

Three-level vision of database

Physical layer, Conceptual (abstraction of the data, relations), Logical layer (further abstraction, )

Real applications use lots and lots of views

# (Virtual) Views

- Defining and Using (Virtual) Views

  - Three-level vision of database

  - Physical layer, Conceptual (abstraction of the data, relations), Logical layer (further abstraction, )

  - Real applications use lots and lots of views

- Benefit of Views

  - Hide some data from some users (Authorization etc.)

  - Make certain queries easier, more natural

  - Modularity of database access

# (Virtual) Views

- Defining and Using (Virtual) Views

  - Three-level vision of database

  - Physical layer, Conceptual (abstraction of the data, relations), Logical layer (further abstraction, )

  - Real applications use lots and lots of views

- Benefit of Views

  - Hide some data from some users (Authorization etc.)

  - Make certain queries easier, more natural

  - Modularity of database access

- Views

  - Query over relations

  - View  $V = \text{ViewQuery}(R_1, R_2, \dots, R_N)$

  - Schema of  $V$  is schema of query result

## Query over views

- V can be thought of as a table

## Query over views

- V can be thought of as a table
- Q is actually rewritten in terms of R1, R2, ....RN

## Query over views

- V can be thought of as a table
- Q is actually rewritten in terms of R1, R2, ....RN
- R1 can also be view

## Query over views

- V can be thought of as a table
- Q is actually rewritten in terms of R1, R2, ....RN
- R1 can also be view
- CREATE View VName (A1, ..., AN) As Query

## Query over views

- V can be thought of as a table
- Q is actually rewritten in terms of R1, R2, ....RN
- R1 can also be view
- CREATE View VName (A1, ..., AN) As Query
- Just define the view and use them (System optimizes in terms of 'base tables')



## Query over views

- V can be thought of as a table
- Q is actually rewritten in terms of R1, R2, ....RN
- R1 can also be view
- CREATE View VName (A1, ..., AN) As Query
- Just define the view and use them (System optimizes in terms of 'base tables')
- View contents not stored

## Query over views

- V can be thought of as a table
- Q is actually rewritten in terms of R1, R2, ....RN
- R1 can also be view
- CREATE View VName (A1, ..., AN) As Query
- Just define the view and use them (System optimizes in terms of 'base tables')
- View contents not stored
- A convenient view is a join

## Modifying Views

- We cannot insert, update, delete as  $V$  isn't stored

## Modifying Views

- We cannot insert, update, delete as  $V$  isn't stored
- To alter  $V$ , we have to modify the base tables

## Materialized Views

- Creates a physical table

# Materialized Views

- Creates a physical table
- But ...

# Materialized Views

- Creates a physical table
- But ...
  - Tables can be very large

# Materialized Views

- Creates a physical table
- But ...
  - Tables can be very large
  - What happens when modifications happen to the base tables



# Materialized Views

- Creates a physical table
- But ...
  - Tables can be very large
  - What happens when modifications happen to the base tables
  - Modification to base data can invalidate the view

## SQLite: Data Types

- NULL – The value is a NULL value

## SQLite: Data Types

- NULL – The value is a NULL value
- INTEGER – a signed integer

## SQLite: Data Types

- NULL – The value is a NULL value
- INTEGER – a signed integer
- REAL – a floating point value

## SQLite: Data Types

- NULL – The value is a NULL value
- INTEGER – a signed integer
- REAL – a floating point value
- TEXT – a text string

## SQLite: Data Types

- NULL – The value is a NULL value
- INTEGER – a signed integer
- REAL – a floating point value
- TEXT – a text string
- BLOB – a blob of data